




Global state

Vladimir Vlassov and Johan Montelius





Global state

Time is very much related to the notion of a *global state*.


If we cannot agree on a time, how should we agree on a global state?

The global state is important:

- Garbage collection
- Deadlock detection
- Termination
- Debugging


Global state

Given a partial order of events, can we say anything about the state of the system?



Local history and Local state

The **history** of a process is a sequence of events: $\langle p_0, p_1, \dots, p_n \rangle$



The **state** of a process is **a description of the process** after (before) an event.

- A state corresponds to a **finite prefix of the process's history**.

Global history and Global state

What is the **global history** of concurrent distributed processes?

- The union of individual histories of all processes?
- Do all unions make sense?


What is the **global state** of a distributed system?

- The union of states of individual processes?

A **global state** corresponds to the initial prefixes of the individual process histories.


Global history and Cut

A **cut** is a subset in the global history up to a specific event in each history.




An event is in the *cut* if it belongs to the events of history up to the specific event.


All cuts are equal, but ...



..some are more equal ..




.. than others



Consistent cuts

For each event e in the cut:

- if f happened before e then
- f is also in the cut.
- In other words, a cut C is **consistent** if, for each event it contains, it also contains all the events that happened before that event.




Consistent global state

A **consistent cut** corresponds to a **consistent global state**.

- It is a **possible state** without contradictions
- it is **consistent with** the **actual execution**
- the actual execution might not have passed through the state, even though it's consistent

Consistent, but not actual states




All real-time cuts are consistent, but who knows the real time?

Linearization


- A **run** is a total ordering of all events in a global history **consistent with each local history**.
- A **linearization** or **consistent run** is a run that describes transitions between **consistent global states**.
- A state S' is **reachable** from state S if there is a **linearization** from S to S' .

Linearization




Possible state transitions

[p0, p1, q0, r0, q1, r1, p2, p3, q2, r2, q3]




Possible state transitions

[q0, p0, p1, r0, q1, r1, p2, p3, q2, r2, q3]



Possible paths



Each path is a consistent run, a linearization, one of which the execution actually took.

Why is this important?

- If we can collect all events and know the happened before order, then we can construct all possible linearizations.
- We know that the actual execution took one of these paths.
- Can we say something about the execution even though we do not know which path was taken?
 - **Yes, we can reason about some property of all the executions, e.g., absence of deadlock, that can be described as a predicate.**

Global state predicate


A global state predicate is a property that is true or false for a global state.

- **Safety** - a predicate is never (or always) true in any state.
- **Liveness** - a predicate that eventually evaluates to true.

How do we determine if a property holds in an execution?

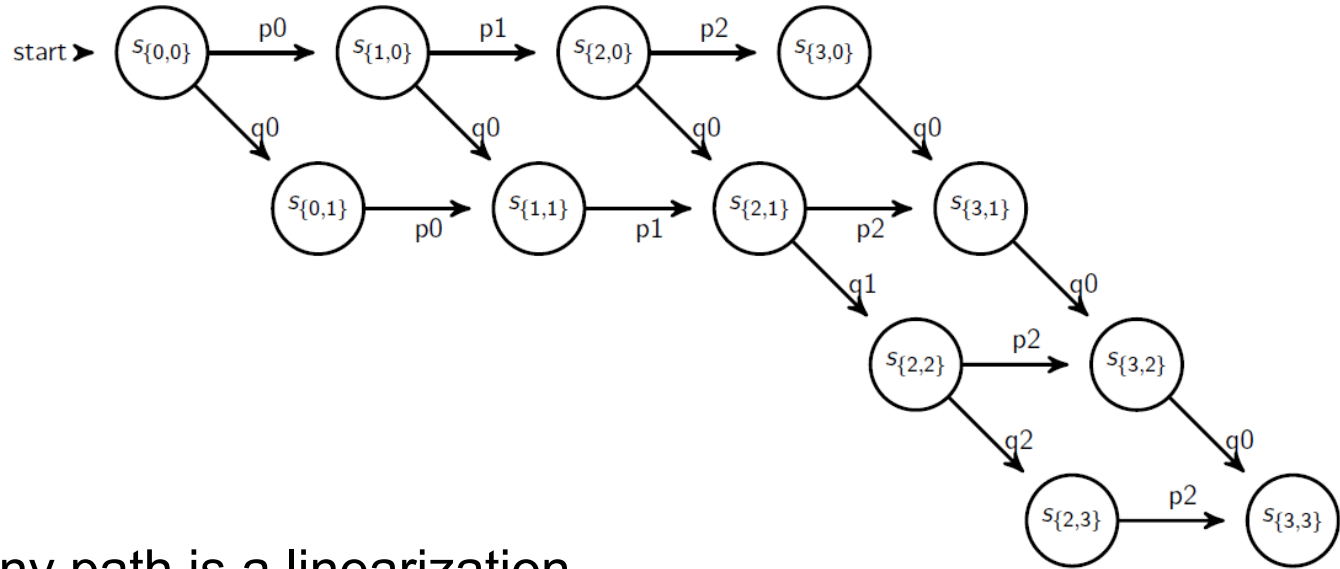
Let's capture all linearizations

Idea - use **vector clocks**, and collect all events of the execution.






An execution lattice



Any path is a linearization.


The actual execution took one path.

Possibly true



If a predicate is true in a consistent global state of the lattice, it is *possibly true* in the execution.

Definitely true



If we cannot find a path from the initial state to the final state without reaching a state for which a predicate is true, then the predicate is *definitely true* during the execution.

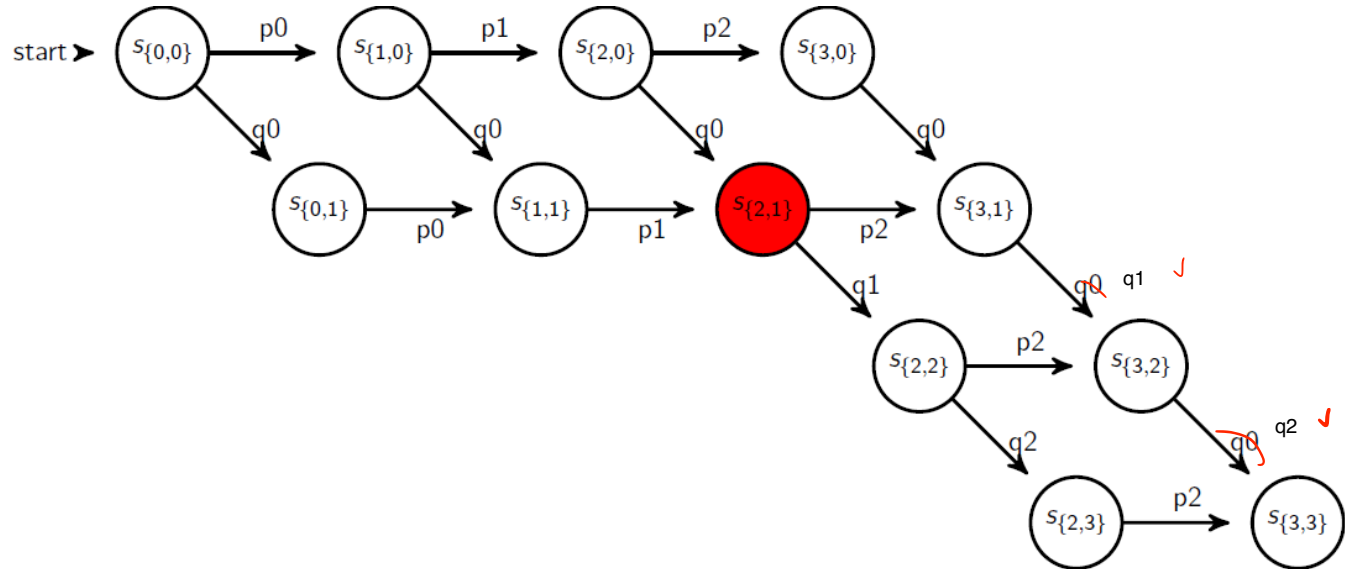
Stable and non-stable

We differentiate between:

- **Stable**: if a predicate is true, it remains true for all reachable states
- **Non-stable**: if a predicate can become true and then later become false

Stable is good

What do I know if a stable predicate is true for state $S_{\{2,1\}}$?





Let's capture a possible state

Idea: capture a consistent global state that was possibly true in the execution.

If a stable predicate is true for this state, it is true in the actual execution.

How do we capture a state?


Snapshot - Chandy and Lamport

A node initiates a snapshot when it receives a *marker*.

- Record the local state and
- send a *marker* on all outgoing channels.
- Record all incoming messages on each channel...
- until you receive a marker.
- When the last channel is closed, you have a local and a set of messages.


Ask one node to initiate the snapshot, collect all local states and messages and construct a global state.

Snapshot markers



What messages are collected by which node?

Snapshot markers




What messages are collected by which node?

What messages are collected by which node?

P: records its state after p1;
channel Q-P: empty;
channel R-P: message r1

Q: records its state after q1;
channel P-Q: message p1;
channel R-Q: empty

R: records its state after r1;
P-R: empty;
Q-R: empty



Snapshot

- Allows us to collect a global state during execution.
- It only allows us to determine stable predicates.



Summary

The happened before order gives us *consistent cuts or consistent global states*.

Using vector clocks, we can time stamp states, *construct all possible linearizations* and evaluate if predicates hold true in the execution.

A snapshot can record a consistent state that can be used to evaluate *stable predicates*.